



4.LOCAL MODEL-AGNOSTIC METHODS: LIME & SHAP



Working of LIME

The Intuition:

- Ignore the original training data and treat the model purely as a black box.
- Create variations of the instance you want to explain and probe the model with these samples.
- Observe how the model's predictions change in response to these variations.
- Fit an interpretable model to the sampled data, weighting samples more heavily when they are closer to the instance of interest.

Mathematical Formulation:

$$\text{explanation}(\mathbf{x}) = \arg \min_{g \in G} L(\hat{f}, g, \pi_{\mathbf{x}}) + \Omega(g)$$

Components:

- \mathbf{x} : Instance to explain
- g : Explanation model (e.g., linear regression, decision tree)
- G : Family of possible explanations
- L : Loss measuring closeness to original model f
- π_x : Proximity measure (defines neighborhood around x)
- $\Omega(g)$: Model complexity penalty (prefer fewer features)



Algorithm Steps

Step 1: Select Instance

Choose instance x for which you want explanation

Step 2: Generate Perturbations

- Create variations of x (method depends on data type)
- Get black box predictions for all perturbed samples

Step 3: Weight by Proximity

- Calculate proximity of each perturbed sample to x
- Closer samples get higher weights
- Uses exponential smoothing kernel: π_x

Step 4: Train Interpretable Model

- Fit weighted model on perturbed dataset
- Common choice: Lasso (for feature selection)
- Select K features: high $\lambda \rightarrow$ decrease λ until K features remain

Step 5: Extract Explanation

- Interpret the local model (coefficients, feature importance, rules)
- This explains the original black box prediction for x



Creating Variations by Data Type

Tabular Data:

- Perturb each feature individually
- Draw from normal distribution with feature's mean and standard deviation
- Issue: Samples from training data center (not around x)
- Issue: Ignores feature correlations → unrealistic data points

Text Data:

- Start with original text
- Randomly remove words
- Representation: Binary features (1 = word included, 0 = removed)
- Each variation = subset of original words

Image Data:

- Segment image into "superpixels" (connected similar-colored pixels)
- Turn superpixels on or off
- Off = replace pixels with user-defined color (e.g., gray)
- Specify probability for turning off each superpixel
- Why not individual pixels? Too many pixels contribute to classes



Strengths:

- Works across tabular, text, AND images (rare)
- Model replacement flexibility (swap models, keep explanation method)
- Short, selective explanations (human-friendly)
- Easy to use: Python (lime), R (iml, DALEX)
- Can use interpretable features different from training features
- Fidelity measure shows reliability of approximation

Critical Limitations:

- Neighborhood choice unsolved (biggest problem—no principled method)
- Instability: Repeating sampling → different explanations
- Unrealistic sampling: Ignores feature correlations
- Can be manipulated: Hide biases (Slack et al. 2020)
- Complexity (K) must be defined in advance



Shapley Values - Game Theory Foundation

The Game Analogy:

- **Game:** The prediction task for a single instance is treated as a cooperative game.
- **Players:** The feature values act as players that work together to produce the final prediction.
- **Gain:** The gain is defined as the model's actual prediction for the instance minus the model's average prediction.
- **Goal:** Determine how to fairly distribute this gain among the feature values.

Shapley Value Definition:

- Average marginal contribution of feature across all possible coalitions

Mathematical Formula:

$$\phi_j(val) = \sum_{S \subseteq \{1, \dots, p\} \setminus \{j\}} \frac{|S|! (p - |S| - 1)!}{p!} (val(S \cup \{j\}) - val(S))$$

Where:

- S = subset of features (coalition without feature j)
- p = total number of features
- val(S) = prediction using only features in S
- Weight term = |S|!(p-|S|-1)!/p! accounts for all possible orderings

$$val_x(S) = \int \hat{f}(x_1, \dots, x_p) d\mathbb{P}_{X_C} - \mathbb{E}[\hat{f}(\mathbf{X})]$$

Where:

X_C = features NOT in coalition S (marginalized out)



Example Apartment Price Prediction

Scenario:

- Model prediction: €300,000
- Average prediction: €310,000
- Difference (gain to distribute): -€10,000

Feature Contributions (Shapley Values):

- park-nearby: +€30,000 (increases value)
- area-50: +€10,000 (increases value)
- floor-2nd: €0 (no effect)
- cat-banned: -€50,000 (decreases value)

Verification:

- Sum: $30,000 + 10,000 + 0 - 50,000 = -€10,000 \checkmark$
- Matches prediction difference (Efficiency property)

Interpretation:

- Cat ban most negatively impacts this apartment
- Park proximity most positively impacts
- All contributions relative to average prediction



Shapley Values - Four Properties

Property 1: Efficiency (Completeness)

$$\sum_{j=1}^p \phi_j = \hat{f}(\mathbf{x}) - \mathbb{E}[\hat{f}(\mathbf{X})]$$

Property 2: Symmetry

- If two features contribute equally to all coalitions, they receive equal Shapley values
- Formally: If $\text{val}(S \cup \{j\}) = \text{val}(S \cup \{k\})$ for all S , then $\phi_j = \phi_k$

Property 3: Dummy

- Features not changing prediction receive Shapley value of zero
- Formally: If $\text{val}(S \cup \{j\}) = \text{val}(S)$ for all S , then $\phi_j = 0$

Property 4: Additivity

- For combined models, sum Shapley values per model
- Useful for ensemble methods (e.g., random forests)

Uniqueness Theorem:

- Shapley values are ONLY solution satisfying all four properties
- Provides solid theoretical foundation



Shapley Estimation Algorithm

Challenge:

- Exact computation requires 2^p coalitions (exponential)
- Infeasible for models with many features

Solution: Monte Carlo Sampling (Štrumbelj & Kononenko, 2014)

Algorithm for Feature j:

Step 1: Draw random instance z from dataset

Step 2: Choose random permutation o of all features

Step 3: Create $x_{\{+j\}}$

- Features appearing before j in permutation $o \rightarrow$ take from x
- Feature $j \rightarrow$ take from z
- Features appearing after $j \rightarrow$ take from z

Step 4: Create $x_{\{-j\}}$

- Same as $x_{\{+j\}}$ but feature j also taken from z

Step 5: Compute marginal contribution

$$\hat{\phi}_j = \frac{1}{M} \sum_{m=1}^M \left(\hat{f}(x_{+j}^{(m)}) - \hat{f}(x_{-j}^{(m)}) \right)$$

Step 6: Repeat M times (e.g., $M = 1000$)

Step 7: Shapley value for j = average over M samples

Note: Repeat entire process for ALL features



Shapley Values Evaluation

Strengths:

- Fair distribution: Efficiency guarantees complete explanation
- Contrastive: Compare to average, subset, or single point
- Solid theory: Only method satisfying four axioms
- Model-agnostic: Works for any model with prediction function

Limitations:

- Computational cost: Requires approximation (2^p coalitions)
- Misinterpretation risk: Not effect of removing from training, but contribution given current features
- Not local in gradient sense: Positive value ≠ increasing feature increases prediction
- Not sparse: Uses ALL features (humans prefer selective explanations)
- Not a model: Returns values only, cannot predict for new inputs
- Unrealistic data points: When features correlated, samples unlikely combinations



What SHAP Adds Beyond Shapley values:

- SHAP introduces new, faster methods for estimating Shapley values, making it more practical for real-world use.
- It provides a unifying theory that connects LIME, Shapley values, and other attribution methods.
- SHAP comes with a rich ecosystem of visualizations, helping you interpret and present explanations more effectively.
- It supports applications not just for tabular data, but also for text and image models.
- With SHAP, there has been a massive increase in popularity and adoption of Shapley-based explanations.
- Its key innovation is representing explanations as a **linear additive model**, where contributions from each feature (the Shapley values) sum up to approximate the model prediction locally

$$g(\mathbf{z}') = \phi_0 + \sum_{j=1}^M \phi_j z'_j$$

Where:

- $\mathbf{z}' \in \{0,1\}^M$ = coalition vector (1 = feature present, 0 = absent)
- $\phi_0 = E[f(X)]$ = base value (average prediction across dataset)
- ϕ_j = SHAP value for feature j (same as Shapley value)
- M = maximum coalition size



SHAP Estimation - KernelSHAP

Step 1: Sample K coalition vectors $z'k \in \{0,1\}^M$ (e.g., K=1000)

Step 2: Convert coalitions to original feature space

- Map $z'k$ to valid data instance using hx
- 1's \rightarrow values from x , 0's \rightarrow values from random instance

Step 3: Get predictions for all coalitions: $f(hx(z'k))$

Step 4: Compute SHAP kernel weights $\pi_x(z'k)$

Step 5: Fit weighted linear regression

- Target: predictions from Step 3
- Features: coalition vectors $z'k$
- Weights: SHAP kernel weights

Step 6: Extract SHAP values from regression coefficients

$$\pi_x(z') = \frac{(M-1)}{\binom{M}{|z'|} |z'| (M - |z'|)}$$

Where:

- $|z'| =$ number of present features (coalition size)
 - Small coalitions (few 1's) \rightarrow large weight
 - Large coalitions (many 1's) \rightarrow large weight
 - Medium coalitions \rightarrow small weight
- Why This Weighting?
- Learn most about features in isolation (small coalitions)
 - Learn about features with all others present (large coalitions)
 - Medium coalitions provide less information about individual contributions



SHAP Estimation - TreeSHAP

Speed Improvement:

- Classic approach: $O(TL \cdot 2^M)$ complexity
- TreeSHAP: $O(TL \cdot D^2)$ complexity
- T = number of trees, L = maximum leaves, D = maximum depth

Why It's Fast:

- Leverages tree structure
- Only considers coalitions that change predictions
- Tree with depth 5 has max 32 distinct predictions (not 2^M coalitions)

Two Versions:

Interventional TreeSHAP (Current Default):

- Computes classic Shapley values
- Uses marginal distribution $P(X_S)$
- Consistent with original Shapley definition

Tree-Path Dependent TreeSHAP:

- Computes conditional SHAP values
- Uses conditional distribution from tree structure
- Was original implementation
- Solves extrapolation but slightly different game

Applicable To:

- Decision trees
- Random forests
- Gradient boosted trees (XGBoost, LightGBM, CatBoost)



SHAP Estimation - Permutation Method

Key Idea:

- Create random feature permutations
- Build coalitions incrementally from both directions
- Reuse predictions across coalition steps

Algorithm:

Step 1: Generate m random permutations of features (e.g., $m=10$)

Step 2: For each permutation $o = (x_o(1), x_o(2), \dots, x_o(M))$:

Forward Pass (Left → Right):

- Coalition 1: \emptyset
- Coalition 2: $\{x_o(1)\}$
- Coalition M : $\{x_o(1), \dots, x_o(M-1)\}$

$$\hat{\phi}_j^{(i)} = \frac{1}{2m} \sum_{k=1}^m (\hat{\Delta}_{o(k),j} + \hat{\Delta}_{-o(k),j})$$

Backward Pass (Right → Left):

- Coalition 1: \emptyset
- Coalition 2: $\{x_o(M)\}$
- Coalition M : $\{x_{-o}(2), \dots, x_{-o}(M)\}$

Step 3: Compute marginal contributions from both passes

Step 4: Average across all permutations:

Visualizations

Force Plot

Visual representation of how features push prediction from baseline to final value

Each feature = force (arrow) pointing up or down

Components:

- Base value (ϕ_0): Starting point = average prediction across dataset
- Red arrows: Features increasing prediction (push right/up)
- Blue arrows: Features decreasing prediction (push left/down)
- Final prediction: Where all forces balance

Example (Penguin Classification):

Base: $P(\text{Adelie}) = 0.5$ (average)



Bill length short (+0.3) → pushes toward Adelie

Body mass high (-0.15) → pushes away from Adelie

Final: $P(\text{Adelie}) = 0.65$

When to Use:

- Explaining single prediction to stakeholders
- Understanding which features dominated decision
- Communicating model behavior intuitively





Strengths:

- Solid theory: Inherits all Shapley value advantages
- Unifies methods: Connects LIME, Shapley, and other approaches
- Fast tree implementation: TreeSHAP key to widespread adoption
- Global interpretations: Feature importance, dependence, interactions, clustering
- Rich visualizations: Force, summary, dependence, interaction plots

Limitations:

- KernelSHAP slow: Impractical for many instances
- Feature dependence: KernelSHAP ignores correlations (like other permutation methods)
- Tree-path dependent issues: Different value function, can give non-zero values to non-influential features
- Can be misinterpreted: Same issues as Shapley values (not removal effect)
- Possible manipulation: Can create misleading explanations to hide biases (Slack et al. 2020)



Aspect	LIME	SHAP
Foundation	Local linear approximation	Game theory (Shapley values)
Weighting	Proximity to instance (exponential kernel)	Coalition size (SHAP kernel)
Theory	No formal guarantees	Four axioms (efficiency, symmetry, dummy, additivity)
Sparsity	Select K features (sparse)	Uses all features
Speed	Fast (one local fit)	Varies: TreeSHAP fast, KernelSHAP slow
Stability	Unstable (different runs vary)	More stable (especially TreeSHAP)
Completeness	Partial explanation	Complete (sum = prediction difference)
Neighborhood	User-defined kernel (problematic)	Defined by Shapley formula
Global use	Limited	Rich ecosystem
Explanation type	Model (can predict)	Values only



Feature Importance

Formula:

$$I_j = \frac{1}{n} \sum_{i=1}^n |\phi_j^{(i)}|$$

Calculation of feature importance:

First, compute the SHAP values $\phi_j^{(i)}$ for **every** instance i in your dataset (with n instances)

- For each feature j , take the absolute value of its SHAP value $|\phi_j^{(i)}|$, so that positive and negative contributions both count.
- Then, average these absolute values over all instances to get the global importance I_j .
- Finally, sort the features in descending order of I_j features with larger average absolute SHAP values are considered more important.

What It Measures:

- Average magnitude of feature's impact on predictions
- NOT direction (uses absolute values)
- Features with high I_j have strong effects (positive or negative)

Example (Penguin Classification):

- Body mass: $I = 0.25$ (changes prediction by 25 percentage points on average)
- Bill length: $I = 0.18$
- Flipper length: $I = 0.12$

Comparison to Other Methods:

- Permutation importance: Measures performance drop when feature shuffled
- SHAP importance: Measures average magnitude of feature attributions
- Both valid, different perspectives

Summary Plot

What It Shows:

- Y-axis: Features (ordered by importance, top = most important)
- X-axis: SHAP value (impact on prediction)
- Each dot: One instance
- Color: Feature value (red = high, blue = low)
- Horizontal spread: Range of SHAP values for that feature


How to Read It:

Example - Body Mass (Top Feature):

- Wide horizontal spread → large range of effects
- Red dots (high body mass) on left (negative SHAP) → decreases P(female)
- Blue dots (low body mass) on right (positive SHAP) → increases P(female)
- Clear pattern: higher value → negative effect

What You Learn:

1. Feature importance (Y-axis ordering)
2. Direction of effects (X-axis: left=decrease, right=increase)
3. Relationship between feature value and effect (color pattern)
4. Variance in effects (horizontal spread)



Dependence Plot

What You Plot:


- X-axis: Feature value ($x_j^{(i)}$)
- Y-axis: SHAP value for that feature ($\phi_j^{(i)}$)
- Each point: One instance from dataset

Mathematical Representation:

- Dataset of points: $\{(x_j^{(1)}, \phi_j^{(1)}), (x_j^{(2)}, \phi_j^{(2)}), \dots, (x_j^{(n)}, \phi_j^{(n)})\}$

Example - Body Mass:

- X-axis: Body mass (3000g to 6000g)
- Y-axis: SHAP value for body mass
- Pattern: Clear negative relationship
- Interpretation: Heavier penguins \rightarrow lower P(female)



Interaction Values

Shapley Interaction Index:

$$\phi_{i,j} = \sum_{S \subseteq \{1, \dots, M\} \setminus \{i, j\}} \frac{|S|! (M - |S| - 2)!}{2(M - 1)!} \delta_{ij}(S)$$

Where:

$$\delta_{ij}(S) = \hat{f}(S \cup \{i, j\}) - \hat{f}(S \cup \{i\}) - \hat{f}(S \cup \{j\}) + \hat{f}(S)$$

What It Measures:

- Pure interaction effect between features i and j
- Effect of having BOTH features beyond their individual contributions
- Positive: Features amplify each other
- Negative: Features cancel each other

Output:

- $M \times M$ matrix per instance
- Diagonal: Main effects (individual SHAP values)
- Off-diagonal: Pairwise interactions

Practical Use:

1. Identify strongest interactions automatically
2. Color dependence plots by interaction feature
3. Understand complex feature relationships